

Introduction To GPU Programming

Martin Schwinzerl, Riccardo de Maria

HSS Section Meeting, CERN

June 3rd, 2020

Goals

- ▶ Overview about the GPU computing landscape (Concepts, Hardware)
 - ▶ Introduce the (currently) two most established frameworks
 - ▶ Provide Simple Examples to get started
 - ▶ Outline the typical workflow within a GPU enhanced Program
 - ▶ Performance Analysis& Constraints
-
- ▶ **Motivate Design-Decisions & Implementation of SIXTRACKLIB**

Note: This is an interactive Jupyter Notebook - all presented examples are designed to work and allow you to experiment with them. Notebook available from: [github.com \(https://github.com/martinschwinzerl/sixtracklib-presentations/tree/master/20200603_hss_section_meeting\)](https://github.com/martinschwinzerl/sixtracklib-presentations/tree/master/20200603_hss_section_meeting)

GPU Hardware Categories

- ▶ **Low-End Gaming: \approx 100 EUR**

- ▷ Still equal or better FP performance than typical CPU in SP and DP
- ▷ DP performance usually much poorer than SP!

- ▶ AMD RX 560 TI 4 GByte

- ▷ 1024 Cores
- ▷ SP Peak Performance 2406 GFLOP/s
- ▷ DP Peak Performance 163 GFLOP/s \approx 1/16 SP

- ▶ NVidia GTX 1050 Ti 4 GByte

- ▷ 768 Cores
- ▷ SP Peak Performance 1981 GFLOP/s
- ▷ DP Peak Performance 62 GFLOP/s = 1/32 SP

- ▶ SP ... IEEE754 32Bit single precision floating point numbers

- ▶ DP ... IEEE754 64Bit double precision floating point numbers

GPU Hardware Categories

- ▶ Low-End Gaming: ≈ 100 EUR
- ▶ **High-End Gaming: ≤ 1000 EUR**
 - ▷ Substantial performance in SP
 - ▷ DP performance varies!
- ▶ NVidia GeForce RTX 2080 8GByte
 - ▷ 3072 Cores
 - ▷ SP Peak Performance 8920 GFLOP/s
 - ▷ DP Peak Performance 279 GFLOP/s = $1/32$ SP
- ▶ AMD RX Vega 64 8 GByte
 - ▷ 3840 Cores
 - ▷ SP Peak Performance 11518 GFLOP/s
 - ▷ DP Peak Performance 720 GFLOP/s = $1/16$ SP
- ▶ SP ... IEEE754 32Bit single precision floating point numbers
- ▶ DP ... IEEE754 64Bit double precision floating point numbers

GPU Hardware Categories

- ▶ Low-End Gaming: ≈ 100 EUR
- ▶ High-End Gaming: ≤ 1 kEUR
- ▶ **Hybrid HPC: ≈ 3 kEUR - 8 kEUR**
 - ▷ Substantial performance in SP
 - ▷ DP performance $\approx 1/2$ SP
- ▶ **Server HPC: ≈ 6 kEUR - 8 kEUR**
 - ▷ Similar performance characteristics as with hybrid HPC, but
 - ▷ No video outputs, i.e. "Accelerator Card"
- ▶ NVidia Titan V 12 GBytes
 - ▷ 3072 Cores
 - ▷ Hybrid card - Video outputs available
 - ▷ SP Peak Performance 12288 GFLOP/s
 - ▷ DP Peak Performance 6144 GFLOP/s = $1/2$ SP
- ▶ AMD Radeon Instinct MI50 16 GBytes
 - ▷ 3840 Cores
 - ▷ SP Peak Performance 13400 GFLOP/s
 - ▷ DP Peak Performance 6700 GFLOP/s = $1/2$ SP
- ▶ SP ... IEEE754 32Bit single precision floating point numbers
- ▶ DP ... IEEE754 64Bit double precision floating point numbers

GPU Hardware Categories

- ▶ Low-End Gaming: ≈ 100 EUR
- ▶ High-End Gaming: ≤ 1 kEUR
- ▶ Hybrid HPC: ≈ 3 kEUR - 8 kEUR
- ▶ Server HPC: ≈ 6 kEUR - 8 kEUR
- ▶ **AI Server HPC: ≈ 6 kEUR - 10 kEUR**

- ▶ Substantial Integer, SP and Half-Precision Performance or
- ▶ TOPS ... Trillion operations per second of dedicated neural network computes
- ▶ Task-specific Hybrid systems with FPGA and/or CPU enhancements
- ▶ No classical GPU, no video outputs
"Accelerator Card"

</div>

- ▶ SP ... IEEE754 32Bit single precision floating point numbers
- ▶ DP ... IEEE754 64Bit double precision floating point numbers

***CPU**s versus **GPU**s*

	CPU	GPU
Cores / Compute Units	2 - 64 (128 HT/SMT)	16 - 80
Arithmetic Units / Core	2-8	64
Peak Performance SP	24-3200 GFLOP/s	1000-19500 GFLOP/s
Peak Performance DP	12-1600 GFLOP/s	50-9700 GFLOP/s

Compared to CPUs, GPUs have:

- More arithmetic units (in particular SP)
- Less logic for control flow
- Less registers per arithmetic units
- Less memory but larger memory bandwidth

</div>

- SP ... IEEE754 32Bit single precision floating point numbers
- DP ... IEEE754 64Bit double precision floating point numbers
- For CPUs: DP performance is roughly 1/2 of SP performance (YMMV)

Introducing Parallelism

- **Task Parallelism:** perform multiple tasks on the same dataset

Example: Check if a number is a prime number by checking divisibility in parallel

Introducing Parallelism

- **Data Parallelism:** perform the same task on multiple datasets

Example: Vector addition

$$\vec{z} = \vec{x} + \vec{y}$$

z_0	=	x_0	+	y_0	←	Data Parallelism: Perform a Single Instruction on Multiple Datasets ⇒ SIMD
z_1	=	x_1	+	y_1	←	
z_2	=	x_2	+	y_2	←	
z_3	=	x_3	+	y_3	←	
\vdots		\vdots		\vdots		
z_i	=	x_i	+	y_i	←	or a
\vdots		\vdots		\vdots		
z_{n-1}	=	x_{n-1}	+	y_{n-1}	←	Single Program on Multiple Datasets ⇒ SPMD

$\vec{x}, \vec{y}, \vec{z} \in \mathbb{R}^n$

Example: Vector Addition on the GPU using CuPy

Cf. <https://github.com/cupy/cupy> (<https://github.com/cupy/cupy>) and <https://cupy.chainer.org/> (<https://cupy.chainer.org/>) for reference

```
In [1]: # If you have a NVidia GPU, one of the simplest ways to offload calculations to  
        the GPU is to use CuPy  
  
import cupy  
import numpy as np  
  
N = 10000000 # 1.0e7 elements  
  
x_host = np.random.rand( N )  
y_host = np.random.rand( N )  
  
# Create vectors of random numbers. CuPy handles all the required steps behind  
the scenes  
x = cupy.asarray( x_host )  
y = cupy.asarray( y_host )  
  
# x and y are cupy entities that "live" on the GPU while  
# x_host and y_host "live" in regular memory  
# The same applies to the result of the calculation z:  
  
z = x + y # Vector addition is actually performed on the GPU
```

```
In [2]: # in order to make use of the result outside of the GPU,  
# we have to convert z into a regular numpy array  
z_host = cupy.asnumpy( z )  
  
# Compare result to calculation on CPU  
print( f"calculation on host and device yield same result: {np.allclose( z_hos  
t, x_host + y_host, rtol=0.0, atol=1e-16)}" )  
  
del x, y, z, x_host, y_host, z_host
```

calculation on host and device yield same result: True

Overview: Frameworks, Libraries, Toolboxes for GPU Programming

- ▶ High-Level Libraries and Applications
 - ▷ CuPy: numpy-like python module for NVIDIA GPUs
 - ▷ Tensorflow: mostly support for NVIDIA GPUs, some opportunities to use other cards via SYCL
 - ▷ Matlab (currently only NVIDIA GPUs are supported)
 - ▷ Julia
 - ▷ Mathematica
 - ▷
- ▶ HPC frameworks & libraries
 - ▷ OpenMP: \geq version 4.5, allows offloading to GPU targets (for AMD, a recent gcc or clang is required)
 - ▷ OpenACC: similar to OpenMP. Traditionally NVIDIA focused but recently got capabilities to offload to AMD
 - ▷ Intel OneAPI
 - ▷ ArrayFire: optimized functions and data-structures to exploit parallelism on both GPUs and CPUs

Overview: Frameworks, Libraries, Toolboxes for GPU Programming

- Low-Level Programming Frameworks

- ROCm: collection of libraries & tools for AMD GPUs, provides **OpenCL** backend
- HIP: abstract language to write GPU programs (AMD, open-source), allows "automated" translation of **CUDA** programs into HIP
- SyCL: "modernized" abstraction on-top of **OpenCL**, uses **SPIR**
- clang + **SPIR-V**: compile code fragments to a intermediate format consumable by modern **OpenCL** implementations
- clang + **PTX**: compile code fragments to an intermediate format consumable by **CUDA** (NVRTC)
- **CUDA** C99/C++1x based framework for NVIDIA GPUs (some parts are open-source, drivers etc. are proprietary)
- **OpenCL**: Vendor neutral computing language, allows to target CPUs, GPUs, FPGAs, Signal Processors

- ▶ Increasingly Important: Code-Transformation tools and intermediate code representations
 - ▷ PTX: Parallel Thread Execution language, used by NVIDIA for CUDA
 - ▷ SPIR, SPIR-V: Standard Portable Intermediate Representation, cross-platform open standard maintained by the Khronos Group; Central in newer graphic standards and newer iterations of OpenCL

Comparison: OpenCL versus CUDA

▸ OpenCL

- ▷ Hardware Neutral: GPUs, CPUs, FPGAs, Signal Processors
- ▷ Common, standardized software stack (Khronos Group) + vendor specific impl.
- ▷ "Programming to the least common denominator"
- ▷ OpenCL 1.2: Still most commonly used, C99 kernel language, limited set of features
- ▷ OpenCL 2.x: some support from AMD and Intel
- ▷ OpenCL 3.0: recently specified, modular refactoring of OpenCL 2.x

▸ CUDA

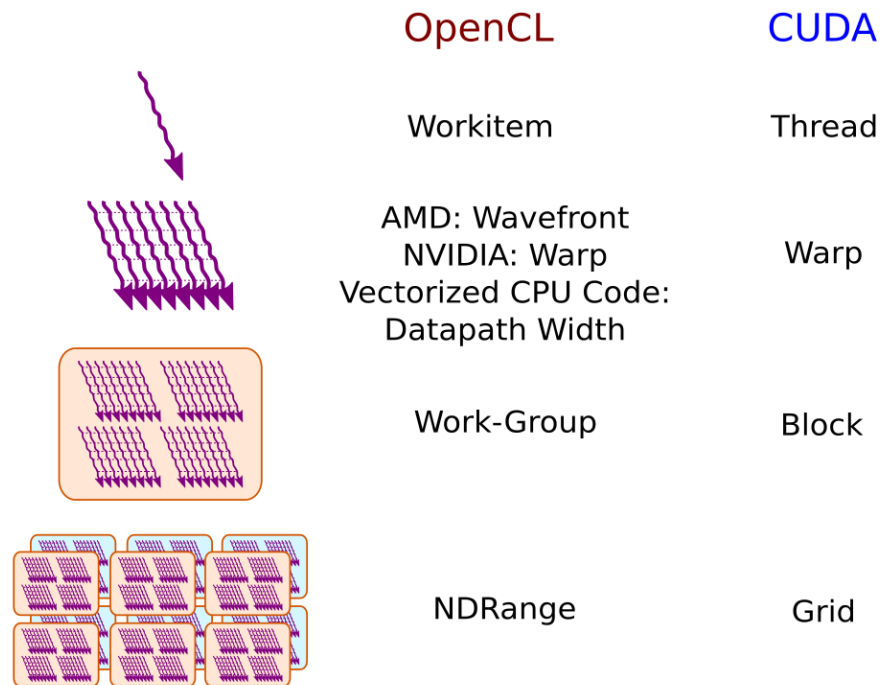
- ▷ Hardware: Only NVIDIA GPUs, no CPU / emulator available
- ▷ Software implementation only available from NVIDIA, includes libraries, debuggers, profilers, memory checkers, etc.
- ▷ Software and hardware tightly integrated (versioning, "compute capabilities")
- ▷ Version 10.x allowing both C99 and C++11/14 kernel language

- ▶ **Primarily:** Run-time compilation of the kernels
- ▶ Kernel code has to be available at run time, compile for selected device at run-time
- ▶ Single-file compilation similar to the default in CUDA available (SyCL)

- ▶ **Primarily:** Single-file compilation. nvcc compiler replaces system compiler for .cu files -> binary compatibility between system compiler, system libraries, graphic stack, etc. required
- ▶ Run-time compilation approach similar to the default with OpenCL is available (NVRTC)

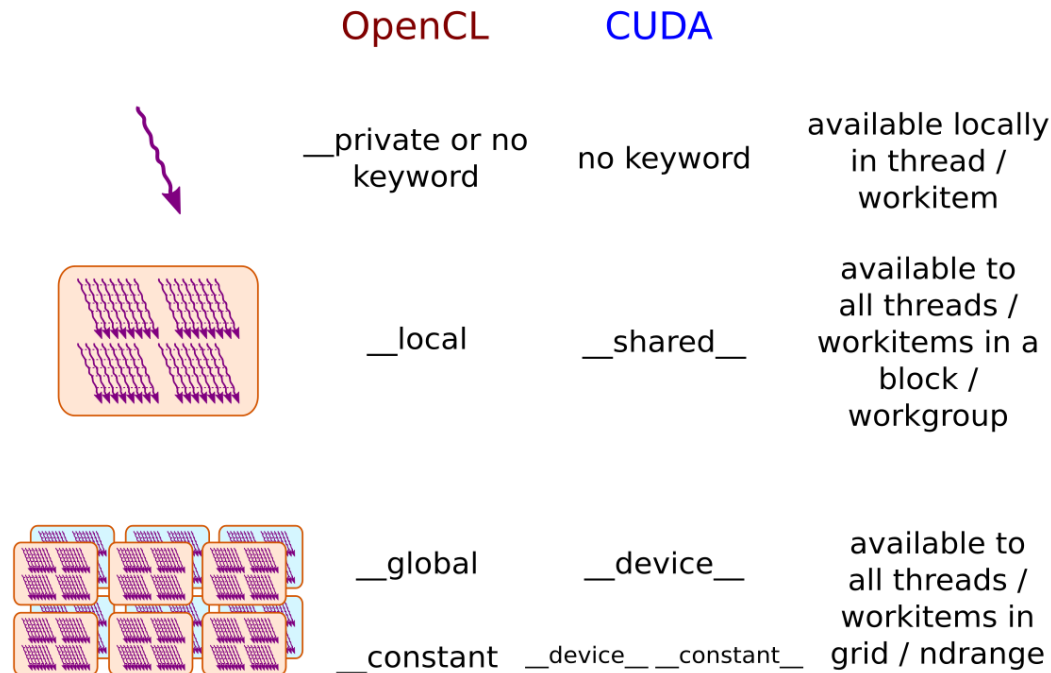
CUDA and OpenCL: Different But Similar - Grid

- Threads are organized in two layers resembling a 3D grid
- Simpler use-cases only requiring 1D or 2D topologies are mapped on the 3D structure




CUDA and OpenCL: Different But Similar - Memory Model

- Hierarchical memory model with different "regions"
- Visibility allows data exchange across individual threads / work-items but may require synchronization
- On CUDA, closely linked to Hardware implementation
- On OpenCL, guarantees about different memory regions are more difficult to give




Structure Of A Simple GPU Program




Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program



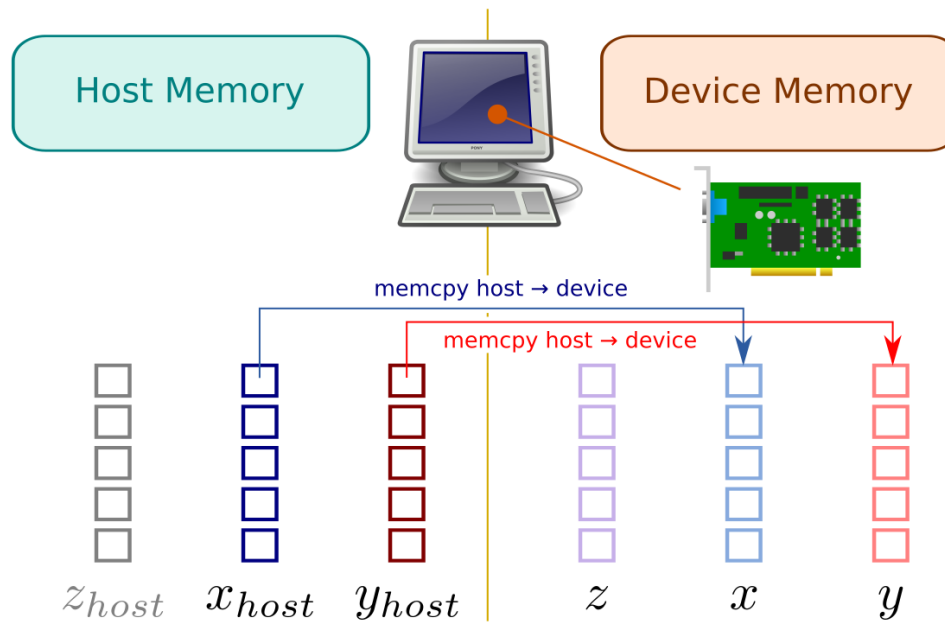
Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program




Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program



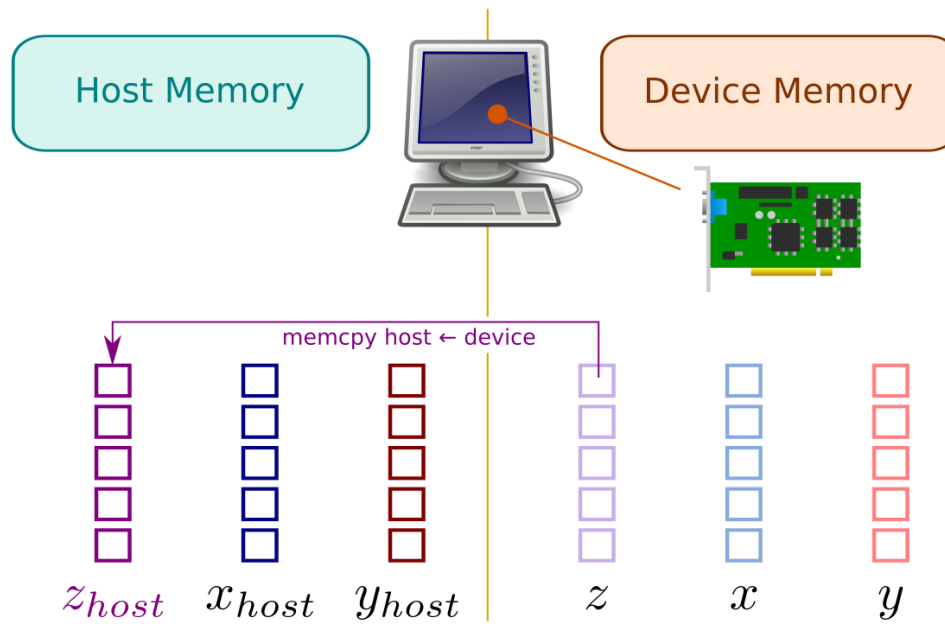
Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program




Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program



Icons: <https://openclipart.org> - License: Public Domain

Structure Of A Simple GPU Program



Icons: <https://openclipart.org> - License: Public Domain

Example: Vector Addition on the GPU using PyCuda

Cf. <https://github.com/inducer/pycuda> (<https://github.com/inducer/pycuda>)
<https://documen.tician.de/pycuda/> (<https://documen.tician.de/pycuda/>) for reference

and

```
In [3]: # Again, a NVIDIA GPU is required for this example to work!
import pycuda.autoinit
import pycuda.driver as cuda
import numpy as np
from pycuda.compiler import SourceModule

# First, we prepare the program that should be applied to our data, i.e. the "Kernel"
vec_add_program = SourceModule(
    """
    __global__ void sum_kernel( double* __restrict__ z,
                                double const* __restrict__ x,
                                double const* __restrict__ y, int const N )
    {
        const int threads_per_block = blockDim.x; /* Assuming 1D grid */
        const int ii = threadIdx.x + blockIdx.x * threads_per_block; /* Assuming 1D
grid */
        if( ii < N )
        {
            z[ ii ] = x[ ii ] + y[ ii ];
        }
    }
    """ )

# Note: even though we are using CUDA, we are able to specify the kernel at run
-time ->
# PyCuda uses the NVRTC implementation!

vec_add_kernel = vec_add_program.get_function( "sum_kernel" )
```

```
In [4]: # then, we prepare the structures on the host side, allocate the  
# required resources on the device side and move things  
  
N = 10000000 # Again, 10^7 elements to the vector  
  
x_host = np.random.randn( N ) # prepare the elements on the host  
y_host = np.random.randn( N )  
  
# allocate the structures on the device  
x = cuda.mem_alloc( x_host.nbytes )  
y = cuda.mem_alloc( y_host.nbytes )  
z = cuda.mem_alloc( max( x_host.nbytes, y_host.nbytes ) )  
  
#transfer the memory from the host to the device; htod ... host to device  
cuda.memcpy_htod( x, x_host )  
cuda.memcpy_htod( y, y_host )
```

```
In [5]: # find the grid dimensions:
threads_per_block = 128 # This is a pessimistic but educated guess :-)
num_blocks = N // threads_per_block
if N % threads_per_block != 0:
    num_blocks += 1

# run the kernel
N_arg = np.int32(N) # pycuda enforces strict type checking
vec_add_kernel( z, x, y, np.int32(N), block=(threads_per_block, 1, 1), grid=(num_blocks,1))

# copy the result back to the host; dtoh ... device to host
z_host = np.empty_like( x_host )
cuda.memcpy_dtoh( z_host, z )

# Compare result to calculation on CPU
print( f"calculation on host and device yield same result: {np.allclose( z_host, x_host + y_host, rtol=0.0, atol=1e-16)}" )

del x, y, z, x_host, y_host, z_host
```

calculation on host and device yield same result: True

Example: Vector Addition on the GPU using PyOpenCL

Cf. <https://github.com/inducer/pyopencl> (<https://github.com/inducer/pycuda>) and <https://document.tician.de/pyopencl/> (<https://document.tician.de/pyopencl/>) for reference

```
In [6]: import pyopencl as cl
import numpy as np

# setup the OpenCL environment - it's a bit more elaborated as with PyCuda:
# Shortcut: ctx = cl.create_some_context() -> create a context with the first a
available device

platforms = cl.get_platforms()
for p in platforms:
    print( p )
```

```
<pyopencl.Platform 'AMD Accelerated Parallel Processing' at 0x7fc00450a250>
<pyopencl.Platform 'NVIDIA CUDA' at 0x3d31de0>
<pyopencl.Platform 'Portable Computing Language' at 0x7fc00dece020>
```

```
In [9]: platform = platforms[ 0 ] # Let's select the first one

# get a list of all available devices for a platform
devices = platform.get_devices()

for d in devices:
    print( d )
```

```
<pyopencl.Device 'gfx900' on 'AMD Accelerated Parallel Processing' at 0x33ec260>
```

```
In [10]: # Again, let's select the first device  
device = devices[ 0 ]  
  
# create a context and a queue  
ctx = cl.Context( [device,] )  
queue = cl.CommandQueue( ctx )
```

```
In [11]: # prepare the program containing the kernel:
vec_add_program = cl.Program( ctx, """
__kernel void sum_kernel( __global double* restrict z,
                          __global double const* restrict x,
                          __global double const* restrict y, int const N )
{
    int const ii = ( int )get_global_id( 0 ); /* Assume 1D Grid, 0 .. 0th element of ndrange */
    if( ii < N )
    {
        z[ ii ] = x[ ii ] + y[ ii ];
    }
}

""" )

# Compile the program containing the kernel
vec_add_program.build()

# get kernel function from vec_add_program
vec_add_kernel = vec_add_program.sum_kernel
```



```
In [12]: # then, we prepare the structures on the host side, allocate the  
# required resources on the device side and move things  
N = 10000000 # Again, 10^7 elements to the vector  
  
# create the structures on the host  
x_host = np.random.rand( N )  
y_host = np.random.rand( N )  
  
# allocate the structures on the device  
x = cl.Buffer( ctx, cl.mem_flags.READ_ONLY, x_host.nbytes )  
y = cl.Buffer( ctx, cl.mem_flags.READ_ONLY, y_host.nbytes )  
z = cl.Buffer( ctx, cl.mem_flags.WRITE_ONLY, max( x_host.nbytes, y_host.nbytes  
) )  
  
#transfer the memory from the host to the device  
cl.enqueue_copy( queue, x, x_host )  
cl.enqueue_copy( queue, y, y_host )
```

```
Out[12]: <pyopencl._cl.NannyEvent at 0x7fc08f85b090>
```

```
In [13]: # find the grid dimensions:
num_work_items = x_host.shape
workgroup_size = None # let the OpenCL impl find the optimal workgroup size

# run the kernel
Narg = np.int32(N) # like with OpenCL
vec_add_kernel( queue, num_work_items, workgroup_size, z, x, y, Narg )

# copy the result back to the host
z_host = np.empty_like( x_host )
cl.enqueue_copy( queue, z_host, z )
```

```
Out[13]: <pyopencl._cl.NannyEvent at 0x7fc08f852360>
```

```
In [14]: # Compare result to calculation on CPU
print( f"calculation on host and device yield same result: {np.allclose( z_hos
t, x_host + y_host, rtol=0.0, atol=1e-16)}" )

del x, y, z, x_host, y_host, z_host
```

```
calculation on host and device yield same result: True
```

Performance Analysis: Scaling

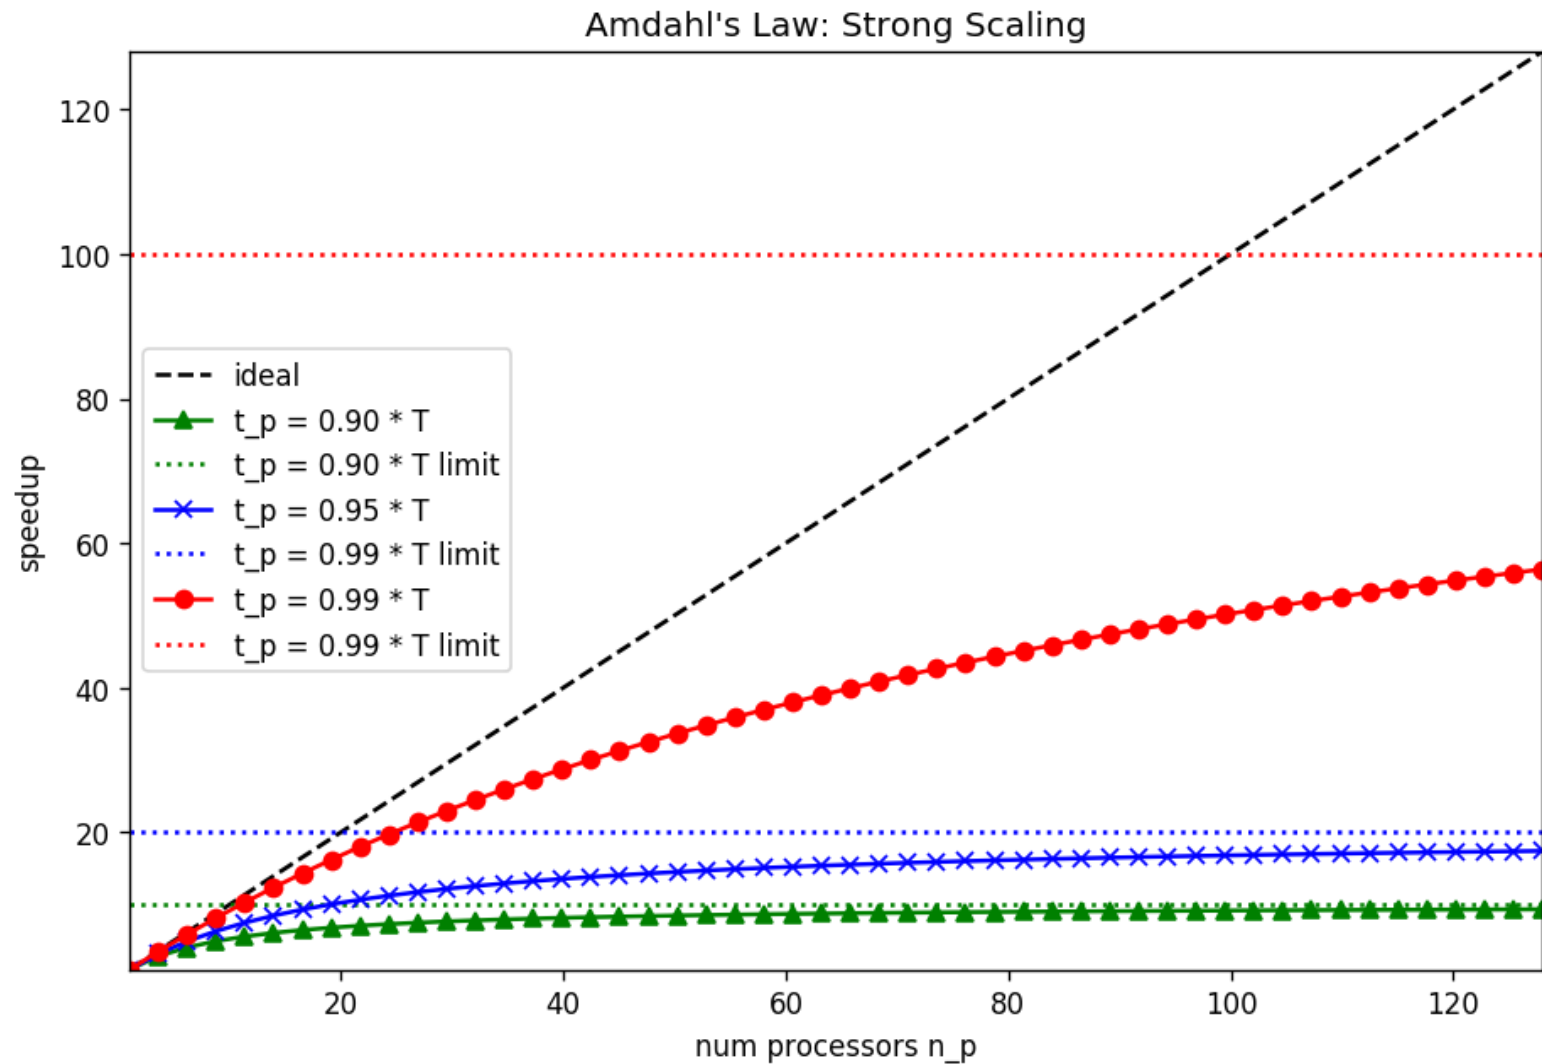
- ▶ Program with a total run-time T
- ▶ Assumption: problem size is **constant** (e.g. add vectors of size N , track N particles, etc.)
- ▶ Question: how much can we speed up the execution of the problem if we parallelise it, i.e. the speed-up $\eta_s(n_p)$

- ▶ $T = t_p + t_s$
- ▶ t_p : fraction of run-time that can be run in "parallel" on n_p "processors"
- ▶ t_s : fraction of run-time that is sequential
- ▶ Amdahls law:

$$\eta_s(n_p) = \frac{T}{t_s + \frac{t_p}{n_p}}$$

Performance Analysis: Scaling

```
In [16]: from helpers import plot_amdahl_scaling  
plot_amdahl_scaling( 1.0, 128.0 )
```



Performance Analysis: Conclusions from Amdahl's Law

- ▶ For a given problem with finite t_s , increases in performance are diminishing with rising n_p
- ▶ Controlling and limiting t_s is crucial to achieve good parallel performance and speedup
- ▶ We assume that $t_s = \text{const.}$ for a given problem size. In practice, $t_s = f(n_p)$
- ▶ Even with t_s very small, scaling to 10^3 or even 10^4 of parallel processes (as in GPUs) under the assumptions of Amdahl is hard

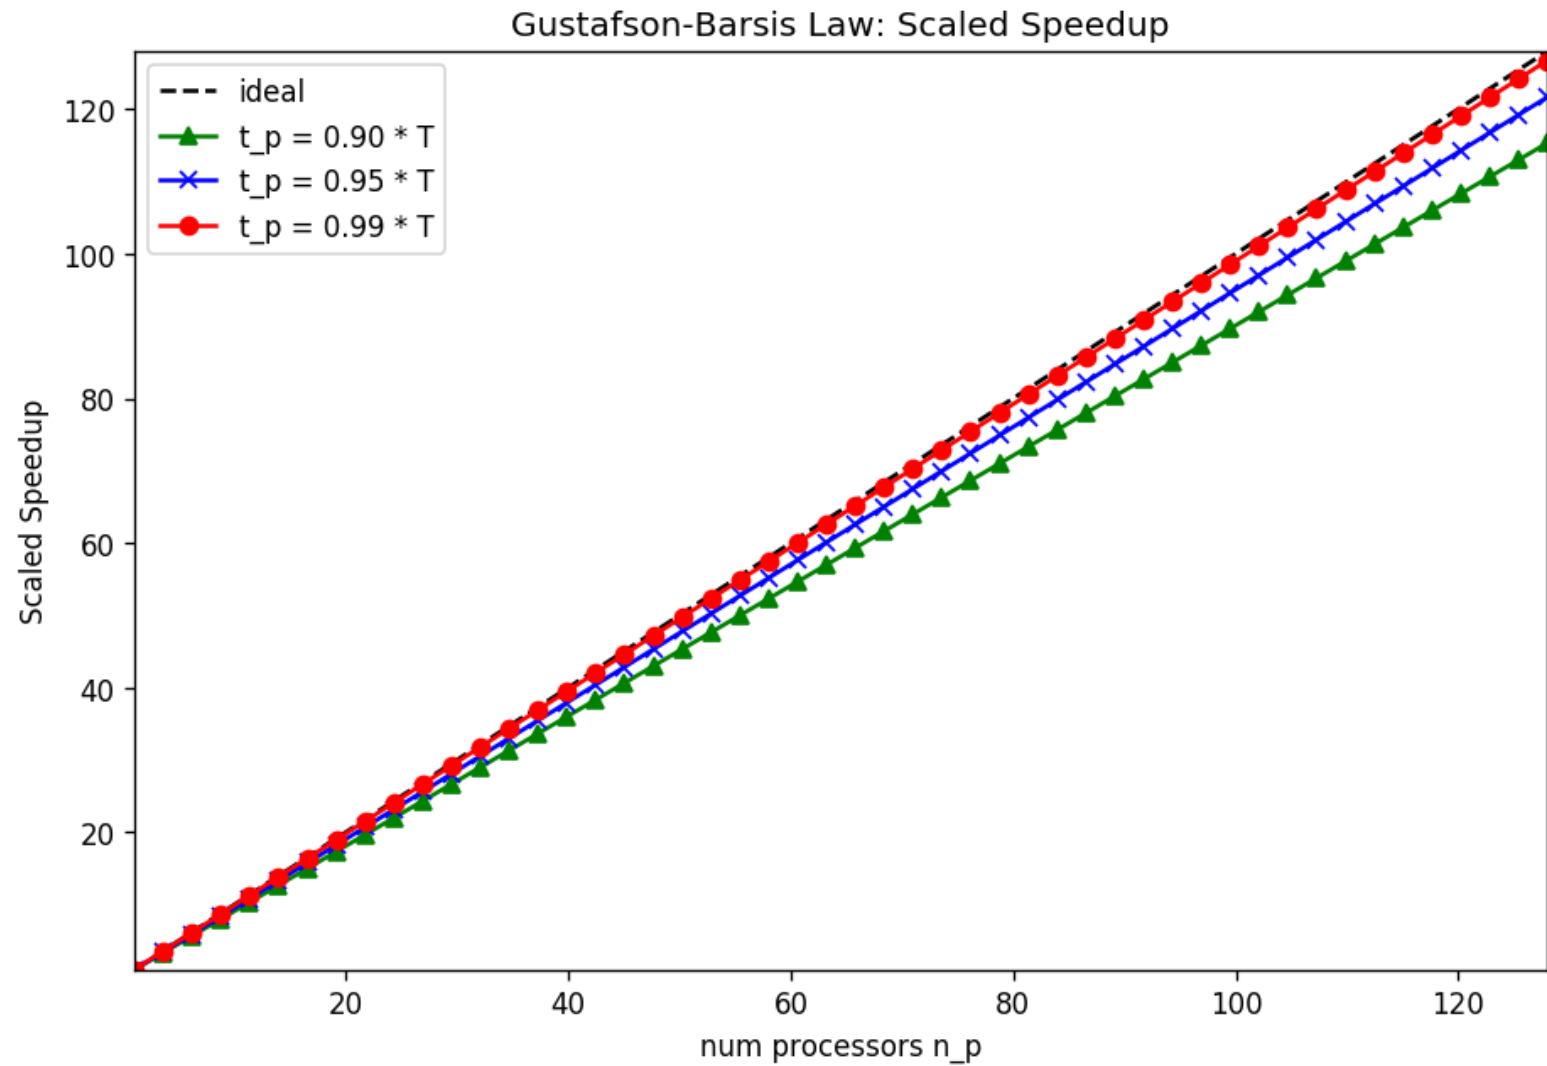
Scaled Speedup, Gustafson-Barsis Law

- ▶ Amdahl's Law is a bit pessimistic ← fixed problem size regardless of n_p
- ▶ What if we can **grow** the problem size together with rising number n_p ?
- ▶ Scaled Speedup:

$$\eta(n_p) = f_s + n_p \cdot f_p$$

- ▶ $f_s = t_s/T$, fraction of the runtime that is serial
- ▶ $f_p = t_p/T$, fraction of the runtime that can be run in parallel on n_p processors.
- ▶ It holds that $f_s + f_p = 1$

```
In [18]: from helpers import plot_gustafson_scaling
plot_gustafson_scaling( 1.0, 128.0 )
```



Contributions to t_s on GPUs

- ▶ Synchronization overhead between "processes" (threads) → Problems that do not require communication between processes are more suitable to be parallelized (ideal: "Embarrassingly Parallel Problem")
 - ▶ Latencies and transfer times for copying memory between host and device
 - ▶ latencies in launching kernels
 - ▶ Latencies for accessing resources (different memory regions!)
 - ▶ During the runtime of the kernel: **thread divergence**
-
- ▶ Thread divergence is a consequence of the SPMD model typically implemented on GPUs
 - ▶ Remember: Threads/work-items are organized in warps/wavefronts
 - ▶ AMD: 64 work-items / wavefront
 - ▶ NVIDIA: 32 threads/workitems / warp
 - ▶ The GPU always executes a multiple of these numbers, even if you only ask to compute 1 work-items / threads
 - ▶ All threads execute the **same program** in lock-step
 - ▶ Dataparallelism: If the program branches depending on the data, code-paths have to be handled sequentially, thus increasing the t_s !

Example: Thread Divergence And Branching

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11      /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13      if( ii < N )
14      {
15          result[ ii ] = a;
16
17          if( x[ ii ] >= 0.0 )
18          {
19              result[ ii ] += x[ ii ];
20          }
21          else
22          {
23              result[ ii ] -= x[ ii ];
24          }
25
26          if( y[ ii ] >= 0.0 )
27          {
28              result[ ii ] += y[ ii ];
29          }
30          else
31          {
32              result[ ii ] -= y[ ii ];
33          }
34
35          result[ ii ] *= c;
36      }
37  }
```

$x \geq 0$



$y \geq 0$



active work-items



Example: Thread Divergence And Branching

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11     /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13     if( ii < N )
14     {
15         result[ ii ] = a;
16
17         if( x[ ii ] >= 0.0 )
18         {
19             result[ ii ] += x[ ii ];
20         }
21         else
22         {
23             result[ ii ] -= x[ ii ];
24         }
25
26         if( y[ ii ] >= 0.0 )
27         {
28             result[ ii ] += y[ ii ];
29         }
30         else
31         {
32             result[ ii ] -= y[ ii ];
33         }
34
35         result[ ii ] *= c;
36     }
37 }
```

$x \geq 0$



$y \geq 0$



active work-items



No Divergence!

Example: Thread Divergence And Branching

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11     /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13     if( ii < N )
14     {
15         result[ ii ] = a;
16
17         if( x[ ii ] >= 0.0 )
18         {
19             result[ ii ] += x[ ii ];
20         }
21         else
22         {
23             result[ ii ] -= x[ ii ];
24         }
25
26         if( y[ ii ] >= 0.0 )
27         {
28             result[ ii ] += y[ ii ];
29         }
30         else
31         {
32             result[ ii ] -= y[ ii ];
33         }
34
35         result[ ii ] *= c;
36     }
37 }
```

$x \geq 0$



$y \geq 0$



active work-items



Divergence

Example: Thread Divergence And Branching

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11     /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13     if( ii < N )
14     {
15         result[ ii ] = a;
16
17         if( x[ ii ] >= 0.0 )
18         {
19             result[ ii ] += x[ ii ];
20         }
21         else
22         {
23             result[ ii ] -= x[ ii ];
24         }
25
26         if( y[ ii ] >= 0.0 )
27         {
28             result[ ii ] += y[ ii ];
29         }
30         else
31         {
32             result[ ii ] -= y[ ii ];
33         }
34
35         result[ ii ] *= c;
36     }
37 }
```

$x \geq 0$



$y \geq 0$



active work-items



Divergence

Example: Thread Divergence And Branching

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11      /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13      if( ii < N )
14      {
15          result[ ii ] = a;
16
17          if( x[ ii ] >= 0.0 )
18          {
19              result[ ii ] += x[ ii ];
20          }
21          else
22          {
23              result[ ii ] -= x[ ii ];
24          }
25
26          if( y[ ii ] >= 0.0 )
27          {
28              result[ ii ] += y[ ii ];
29          }
30          else
31          {
32              result[ ii ] -= y[ ii ];
33          }
34
35          result[ ii ] *= c;
36      }
37  }
```

$x \geq 0$



$y \geq 0$



Assumption: already
Converged again

active work-items



Example: Alternative, Branchless Formulation of Kernel

```
1  /* Example: Naive implemenation of f(c,a,x,y) = c * ( a + abs(x) + abs(y) ) */
2
3  __kernel void c_times_a_plus_abs_x_plus_y(
4      __global double* restrict result,
5      double const c, double const a,
6      __global double const* restrict y,
7      __global double const* restrict x, int const N )
8  {
9      int const ii = ( int )get_global_id( 0 );
10
11      /* result[ ii ] = c * ( a + | x[ ii ] | + | y[ ii ] | ) */
12
13      if( ii < N )
14      {
15          double temp = a;
16          temp += x[ ii ] * ( double )( ( x[ ii ] > 0.0 ) - ( x[ ii ] < 0.0 ) );
17          temp += y[ ii ] * ( double )( ( y[ ii ] > 0.0 ) - ( y[ ii ] < 0.0 ) );
18          temp *= c;
19
20          result[ ii ] = temp;
21      }
22 }
```

Summary & Conclusions

- ▶ GPUs offer a large number of threads and significant floating-point computing power
- ▶ There are many options to get started with GPU programming without starting to go fully towards low-level programming
- ▶ The two current main frameworks, OpenCL and CUDA, differ in scope and implementation but are similar enough that a common implementation of Kernels is possible
- ▶ Controlling and Limiting the sequential run-time component is crucial to achieve good scaling and parallel performance
- ▶ Thread Divergence can be a key contributor to the sequential runtime-portion
- ▶ Not covered: optimal grid dimensions and hardware utilization

**Thank You For Your
Attention!**